

Melanoma_Project2

Grecia De La Cerda

4/30/2020

Introduction

Since it's already summer in Texas I thought a good dataset to look at would be the skin's worse enemy, melanomas. The dataset I chose was "melanoma" in the R package boot. It contains 205 observations with prognostic variables after patients had a complete removal of their tumors. The variables measure the following: "time"- survival time in days post-op "status"-end of study livelihood status (1=dead,2=alive,3=death unrelated) "sex"-patients sex (1=M / 0=F) "age"- age at the time of surgery "thickness"- tumor thickness (mm) "ulcer"- ulceration indications (1=present,0=absent) Additionally the project asks for a categorical variable with 2-5 groupings so I took the age and separated into the quarters/stages of life (0-100 broken into four quarters) and added the variable in the data frame.

```
#Dataset
library(tidyverse)
library(boot)
data(melanoma)
melanoma<-as.data.frame(melanoma)
attach(melanoma)
quarter<-cut(age ,breaks = c(0,25,50,75,100), labels = c("first","second","third","fourth"), right=FALSE)
melanoma$quarter<-quarter
```

MANOVA

I ran a MANOVA to test if time or tumor thickness differed by life stage ("quarter") and it came back with significant results (p-value= 0.00469). I then did a univariate ANOVA from my MANOVA object to find that the survival time in days, after the tumor removal operation, differs by life stage.

There was a total of 6 tests done (1 MANOVA, 2 ANOVAs, 3 post hoc t tests on significant data) and a calculated type I error equating to 0.265, roughly 27%. I then conducted a bonferroni correction to get a cut off of 0.008.

Assumptions for MANOVA were not met because the data rejected the null hypothesis that the DVs (time, thickness) are equal. For the ANOVA, since the variable 'time' was significant, the null hypothesis was also rejected of equal groupings. Then running the t test comparing time against different quarters, significant difference was found between the following quarters: - first quarter vs. fourth quarter - second quarter vs. third quarter - second quarter vs. fourth quarter

```
#Manova code
```

```
man1<-manova(cbind(time,thickness)~quarter)
summary(man1)
```

```
##           Df    Pillai approx F num Df den Df  Pr(>F)
## quarter      3 0.090497    3.1753      6   402 0.00469 **
## Residuals 201
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary.aov(man1)
```

```
## Response time :
##           Df      Sum Sq Mean Sq F value    Pr(>F)
## quarter      3  20639724 6879908   5.8546 0.0007481 ***
## Residuals   201 236200385 1175126
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Response thickness :
##           Df      Sum Sq Mean Sq F value    Pr(>F)
## quarter      3    46.83   15.611   1.8035 0.1478
## Residuals   201 1739.85    8.656
```

```
pairwise.t.test(time,quarter, p.adj="none")
```

```
##
## Pairwise comparisons using t tests with pooled SD
##
## data:  time and quarter
##
##           first    second    third
## second 0.62240 -          -
## third  0.06170 0.00676 -
## fourth 0.00361 0.00055 0.04091
##
## P value adjustment method: none
```

```
1-.956
```

```
## [1] 0.2649081
```

```
.05/6
```

```
## [1] 0.008333333
```

Randomization Test and Plot

Much like the MANOVA, using the PERMANOVA randomization test, it confirms my findings in the last chunk of code. Once again, we reject the null hypothesis of equal euclidean distances per group.

```
#Randomization test code
```

```
library(vegan)
dists<-melanoma%>%select(time, thickness)%>%dist()
adonis(dists~quarter,data=melanoma)
```

```
##
## Call:
## adonis(formula = dists ~ quarter, data = melanoma)
##
## Permutation: free
## Number of permutations: 999
##
## Terms added sequentially (first to last)
##
```

```
##           Df SumsOfSqs MeanSqs F.Model      R2 Pr(>F)
## quarter      3  20639771 6879924  5.8546 0.08036 0.002 **
## Residuals 201 236202125 1175135          0.91964
## Total      204 256841895          1.00000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
#plot
F_obs<-5.8546

SST<-sum(dists^2)/205

Fs<-replicate(1000,{
  new<-melanoma%>%mutate(quarter=sample(quarter))

  SSW<-new%>%group_by(quarter)%>%select(time,thickness)%>%do(d=dist(. [2:3], "euclidean"))%>%ungroup()%>%
  ((SST-SSW)/3)/(SSW/201)
})

{hist(Fs,prob=T);abline(v=F_obs, col="red", add=T)}
```



```
mean(Fs>F_obs)
```

```
## [1] 0
```

Linear Regression

The coefficients after mean centering the data lead to the creation of this following model: $ulcer = 0.400 + 0.062(thickness_c) + 0.092(sex) + 0.014(thickness_c:sex)$. The proportion of the variation in the outcome of my model is the adjusted R squared value from the mean centered fitted model of 0.178.

```
#Linear regression code
library(sandwich)
library(lmtest)
fit<-lm(ulcer~ thickness*sex, data=melanoma)
summary(fit)

##
## Call:
## lm(formula = ulcer ~ thickness * sex, data = melanoma)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3828 -0.3196 -0.2578  0.4611  0.7721
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.21788    0.05423   4.018 8.30e-05 ***
## thickness     0.06239    0.01465   4.260 3.14e-05 ***
## sex           0.05236    0.09451   0.554  0.580
## thickness:sex 0.01350    0.02183   0.619  0.537
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4511 on 201 degrees of freedom
## Multiple R-squared:  0.19, Adjusted R-squared:  0.1779
## F-statistic: 15.71 on 3 and 201 DF, p-value: 3.222e-09

coef(fit)

##      (Intercept)      thickness          sex thickness:sex
##      0.21788331      0.06239172      0.05236220      0.01350075

mean(thickness)

## [1] 2.919854

data.frame(thickness_c=melanoma$thickness-mean(melanoma$thickness))

##      thickness_c
## 1      3.84014634
## 2     -2.26985366
## 3     -1.57985366
## 4     -0.01985366
## 5      9.16014634
## 6      1.92014634
## 7      2.24014634
## 8      0.30014634
## 9      9.96014634
## 10     4.49014634
## 11     1.27014634
## 12    -2.75985366
```

13 0.95014634
14 1.92014634
15 -0.49985366
16 9.64014634
17 2.88014634
18 4.14014634
19 2.56014634
20 4.81014634
21 10.93014634
22 -0.57985366
23 1.27014634
24 1.12014634
25 1.92014634
26 -2.59985366
27 5.62014634
28 -0.33985366
29 0.64014634
30 0.62014634
31 -1.94985366
32 1.91014634
33 -1.29985366
34 3.52014634
35 11.74014634
36 -0.33985366
37 0.95014634
38 0.62014634
39 -1.57985366
40 -0.67985366
41 0.95014634
42 0.62014634
43 14.50014634
44 -1.62985366
45 0.30014634
46 -1.62985366
47 1.59014634
48 5.46014634
49 -0.97985366
50 -2.75985366
51 -0.33985366
52 -1.62985366
53 -2.75985366
54 -1.29985366
55 -1.62985366
56 -0.81985366
57 -2.59985366
58 -2.10985366
59 -1.78985366
60 2.24014634
61 -1.29985366
62 -1.54985366
63 -2.67985366
64 -2.10985366
65 -1.62985366
66 -1.62985366

67 -1.94985366
68 -1.78985366
69 2.88014634
70 -1.62985366
71 -2.43985366
72 -1.29985366
73 -0.65985366
74 -2.33985366
75 -1.94985366
76 -0.33985366
77 -2.10985366
78 0.62014634
79 -1.94985366
80 -1.13985366
81 -0.97985366
82 -1.62985366
83 0.30014634
84 -1.38985366
85 -1.62985366
86 -1.29985366
87 -1.29985366
88 -2.59985366
89 1.92014634
90 -1.62985366
91 -1.94985366
92 0.14014634
93 0.62014634
94 -1.29985366
95 -0.33985366
96 -0.97985366
97 -2.10985366
98 4.81014634
99 -1.94985366
100 9.96014634
101 -0.33985366
102 1.17014634
103 -2.27985366
104 -1.94985366
105 0.30014634
106 -1.29985366
107 0.95014634
108 -2.59985366
109 -2.59985366
110 0.30014634
111 -0.65985366
112 0.14014634
113 -0.33985366
114 -2.26985366
115 -1.78985366
116 -2.10985366
117 -1.94985366
118 -1.15985366
119 -0.97985366
120 -2.26985366

121 -1.94985366
122 2.72014634
123 6.74014634
124 -2.81985366
125 2.56014634
126 -0.65985366
127 1.91014634
128 -1.94985366
129 -1.94985366
130 2.24014634
131 -2.10985366
132 -0.01985366
133 0.95014634
134 -0.97985366
135 -2.75985366
136 -2.27985366
137 -0.65985366
138 -1.46985366
139 1.90014634
140 -1.62985366
141 4.97014634
142 -2.10985366
143 0.62014634
144 -1.62985366
145 -2.27985366
146 0.30014634
147 -1.46985366
148 -2.43985366
149 -0.97985366
150 -2.75985366
151 -2.75985366
152 -1.62985366
153 -0.97985366
154 0.62014634
155 -2.10985366
156 -2.26985366
157 4.17014634
158 -2.75985366
159 -1.29985366
160 -1.29985366
161 -1.62985366
162 3.20014634
163 -2.43985366
164 -2.27985366
165 0.30014634
166 -0.97985366
167 -0.33985366
168 -0.33985366
169 -2.10985366
170 -2.10985366
171 0.30014634
172 -2.59985366
173 0.30014634
174 -0.17985366

```
## 175 1.92014634
## 176 -1.29985366
## 177 -2.26985366
## 178 -1.46985366
## 179 -2.26985366
## 180 -1.62985366
## 181 -1.29985366
## 182 0.62014634
## 183 0.30014634
## 184 -2.26985366
## 185 -1.88985366
## 186 4.17014634
## 187 -1.62985366
## 188 -2.26985366
## 189 -1.13985366
## 190 9.32014634
## 191 5.14014634
## 192 -2.10985366
## 193 -0.81985366
## 194 0.95014634
## 195 -2.26985366
## 196 -0.97985366
## 197 -2.26985366
## 198 -0.81985366
## 199 -0.97985366
## 200 -1.78985366
## 201 4.14014634
## 202 3.20014634
## 203 -2.43985366
## 204 -0.65985366
## 205 -0.01985366
```


```
melanoma$thickness_c<-melanoma$thickness-mean(melanoma$thickness)
fit1<-lm(ulcer~thickness_c*sex, data=melanoma)
summary(fit1)
```

```
##
## Call:
## lm(formula = ulcer ~ thickness_c * sex, data = melanoma)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3828 -0.3196 -0.2578  0.4611  0.7721
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.40006    0.04068   9.833 < 2e-16 ***
## thickness_c     0.06239    0.01465   4.260 3.14e-05 ***
## sex            0.09178    0.06600   1.391  0.166
## thickness_c:sex 0.01350    0.02183   0.619  0.537
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4511 on 201 degrees of freedom
## Multiple R-squared:  0.19, Adjusted R-squared:  0.1779
```




```
## F-statistic: 15.71 on 3 and 201 DF, p-value: 3.222e-09
```

```
qplot(x=thickness_c, y=ulcer, color= sex, data= melanoma)+  
  stat_smooth(method = "lm", se= FALSE, fullrange = TRUE)
```



```
resids<-fit1$residuals  
fitvals<-fit1$fitted.values  
ggplot()+geom_point(aes(fitvals,resids))+geom_hline(yintercept = 0, col="red")
```



```
bptest(fit1)
```

```
##
## studentized Breusch-Pagan test
##
## data: fit1
## BP = 10.826, df = 3, p-value = 0.0127
```

```
ks.test(resids, "pnorm", sd=sd(resids))
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: resids
## D = 0.26674, p-value = 4.282e-13
## alternative hypothesis: two-sided
```

```
coeftest(fit1, vcov. = vcovHC(fit1))[1:2]
```

```
##              Estimate Std. Error
## (Intercept)  0.40005801 0.04312965
## thickness_c  0.06239172 0.01748523
## sex          0.09178241 0.06542109
## thickness_c:sex 0.01350075 0.03141322
```

Bootstrapped SE

```
#Bootstrapped SE, CIs Code
boot_dat<-sample_frac(melanoma,replace = T)
```

```
samp_distn<-replicate(5000, {
  boot_dat<-sample_frac(melanoma, replace= T)
  fit<-lm(ulcer~thickness_c*sex,data=melanoma)
  coef(fit)
})

samp_distn%>%t%>%as.data.frame%>%gather%>%group_by(key)%>%summarize(lower=quantile(value,.025),upper=quantile(value,.975))

## # A tibble: 4 x 3
##   key                lower upper
##   <chr>              <dbl> <dbl>
## 1 (Intercept)      0.400  0.400
## 2 sex              0.0918 0.0918
## 3 thickness_c      0.0624 0.0624
## 4 thickness_c:sex  0.0135 0.0135
```

Logistic Regression

The log(odds) model from the coefficients looks as follows: $\log(\text{odds}) = -1.708 + 0.448(\text{thickness}) + 0.0048(\text{age})$. Holding age constant, going up one unit of tumor thickness increases the logit score by 0.448. The accuracy rate is 0.7560976. The TPR or the probability of detecting ulceration correctly is $53/90 = 0.5888889$. The TNR rate is $102/115 = 0.8869565$. The PPV rate is $53/66 = 0.8030303$. For the calculated AUC the value is 0.820, making it a good enough model. After doing a CV, looking the average out of sample Accuracy, Sensitivity, and Recall, the model looks a slightly worse out of the original sample.

#Logistic Regression Code

```
fit2<-glm(ulcer~thickness + age, data= melanoma, family = "binomial")
coeftest(fit2)
```

```
##
## z test of coefficients:
##
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.7082538  0.5445886 -3.1368  0.001708 **
## thickness    0.4482084  0.0878074  5.1045 3.318e-07 ***
## age          0.0048853  0.0096873  0.5043  0.614051
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coef(fit2)%>%round(5)%>%data.frame
```

```
##
## (Intercept) -1.70825
## thickness    0.44821
## age          0.00489
probs<- predict(fit2,type="response")
table(predict=as.numeric(probs>.5), truth=melanoma$ulcer)%>%addmargins
```

```
##           truth
## predict    0    1 Sum
##      0    102  37 139
##      1     13  53  66
##      Sum   115  90 205
```

```

melanoma$logit<-predict(fit2,type="link")
(102+53)/205

## [1] 0.7560976
53/90

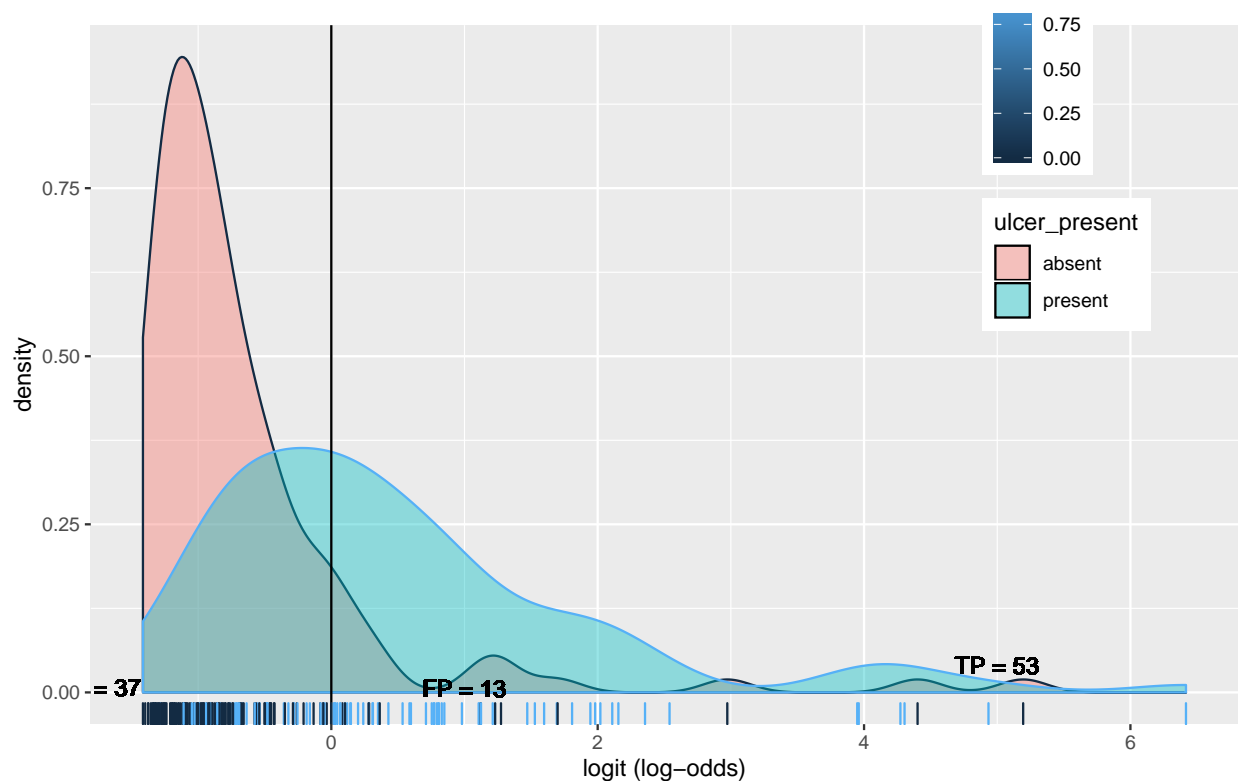
## [1] 0.5888889
102/115

## [1] 0.8869565
53/66

## [1] 0.8030303
data<-melanoma%>%
  mutate(ulcer_present= case_when(
    ulcer==1 ~ "present", ulcer==0 ~ "absent"))
melanoma<-data

melanoma%>%ggplot()+geom_density(aes(logit,color=ulcer,fill=ulcer_present), alpha=.4)+
  theme(legend.position=c(.85,.85))+geom_vline(xintercept=0)+xlab("logit (log-odds)")+
  geom_rug(aes(logit,color=ulcer))+
  geom_text(x=-5,y=.07,label="TN = 102")+
  geom_text(x=-1.75,y=.008,label="FN = 37")+
  geom_text(x=1,y=.006,label="FP = 13")+
  geom_text(x=5,y=.04,label="TP = 53")

```




```

#ROC and AUC
library(plotROC)

```

```
ROCplot<-ggplot(melanoma)+geom_roc(aes(d=ulcer_present,m=probs), n.cuts=0)
ROCplot
```



```
calc_auc(ROCplot)
```

```
## PANEL group AUC
## 1 1 -1 0.8201449
```

```
#extra
```

```
class_diag <- function(probs,truth){
```

```
#CONFUSION MATRIX: CALCULATE ACCURACY, TPR, TNR, PPV
```

```
tab<-table(factor(probs>.5,levels=c("FALSE","TRUE")),truth)
```

```
acc=sum(diag(tab))/sum(tab)
```

```
sens=tab[2,2]/colSums(tab)[2]
```

```
spec=tab[1,1]/colSums(tab)[1]
```

```
ppv=tab[2,2]/rowSums(tab)[2]
```

```
if(is.numeric(truth)==FALSE & is.logical(truth)==FALSE) truth<-as.numeric(truth)-1
```

```
#CALCULATE EXACT AUC
```

```
ord<-order(probs, decreasing=TRUE)
```

```
probs <- probs[ord]; truth <- truth[ord]
```

```
TPR=cumsum(truth)/max(1,sum(truth))
```

```
FPR=cumsum(!truth)/max(1,sum(!truth))
```

```
dup<-c(probs[-1]>=probs[-length(probs)], FALSE)
```

```
TPR<-c(0,TPR[!dup],1); FPR<-c(0,FPR[!dup],1)
```

```

n <- length(TPR)
auc<- sum( ((TPR[-1]+TPR[-n])/2) * (FPR[-1]-FPR[-n]) )

data.frame(acc,sens,spec,ppv, auc)
}
class_diag(probs,ulcer)

```

```

##          acc          sens          spec          ppv          auc
## 1 0.7560976 0.5888889 0.8869565 0.8030303 0.8201449

```

```

library(pROC)
auc(ulcer,probs)

```

```

## Area under the curve: 0.8201

```

```

#CV
set.seed(1234)
k= 10
data<-melanoma[sample(nrow(melanoma)),]
folds<-cut(seq(1:nrow(melanoma)),breaks=k,labels=F)
diags<-NULL
for(i in 1:k){
train<-data[folds!=i,]
test<-data[folds==i,]
truth<-test$ulcer
fit<-glm(ulcer~thickness+age,data=train,family="binomial")
probs<-predict(fit,newdata = test,type="response")

diags<-rbind(diags,class_diag(probs,truth))
}

summarize_all(diags,mean)

```

```

##          acc          sens          spec          ppv          auc
## 1 0.7509524 0.569725 0.8969048 0.7992857 0.8100212

```

LASSO

The variables that were retained and used for the 10-fold CV at the end (that did offer a better fit) was status, age, and ulcer for time of survival post operation.

```

#LASSO Code

```